

Node

label: myKey=label3



POD

Kubernetes Preferred Node Affinity: Soft Constraints

Flexible pod placement with weighted preferences

Node

label: myKey=label1

- ```
affinity:
 nodeAffinity:
 requiredDuringSchedulingIgnoredDuringExecution:
 - labelKey:
 operator: In
 values:
 - label1
 - label2
```

# Understanding Preferred Node Affinity

## ☰ Key Concepts

- ✓ **preferredDuringSchedulingIgnoredDuringExecution** - Soft rule (scheduler tries, but won't block)
- ✓ **weight** - How much to favor this rule (1 = low, 100 = high)
- ✓ **preference** - What node attributes are preferred

## ⚖ Weight Range

Use higher weights for stronger preferences

● 1-30: Low ● 31-70: Medium ● 71-100: High

## YAML Example

```
affinity: nodeAffinity:
 preferredDuringSchedulingIgnoredDuringExecution:
 - weight: 1 # ← Importance (1-100)
 preference:
 matchExpressions:
 - key: size
 operator: In
 values: - small
```

## ↔ Behavior Scenarios

| Scenario                               | Outcome                                          |
|----------------------------------------|--------------------------------------------------|
| size=small node exists + has resources | <b>Scheduled there</b> (preferred)               |
| No size=small node                     | <b>Scheduled elsewhere</b> (no failure!)         |
| Multiple preferred rules               | <b>Scores nodes</b> → picks highest total weight |

## ⤵ Key Benefits

- ✓ **Flexible scheduling** - Pods always schedule somewhere
- ✓ **Optimization focus** - Best-effort placement for better performance
- ✓ **Cost-effective** - Can prefer cheaper instances when available
- ✓ **Graceful fallback** - Works even if preferred nodes unavailable

# Lab: Preferred Affinity in Action

## 1 ➤ Label Node as "small"

```
List nodes
kubectl get nodes

Label ONLY ONE node
kubectl label node k3s-node1 size=small

Verify
kubectl get nodes --show-labels | grep -E
"k3s-node1|k3s-node2"
```

## 2 🚤 Deploy Pod

```
Create namespace
kubectl create namespace learning

Apply Pod with preferred affinity
kubectl apply -f pod-with-preferred-node-
affinity.yml

Check where it runs
kubectl get pods -n learning -o wide
```

## 3 ✓ Expected Result

```
Pod runs on k3s-node1 (the "small" node)
But if k3s-node1 is full, it might run on
k3s-node2
→ That's OK with preferred affinity!
```

## 4 ⚠️ Test Fallback Behavior

```
Taint "small" node to make it unschedulable
kubectl taint node k3s-node1
test=unschedulable:NoSchedule

Deploy a SECOND Pod (same spec)
kubectl run nnappone2 -n learning --image=nginx --
restart=Never \\ --overrides='{"spec": {"affinity":
{"nodeAffinity":
{"preferredDuringSchedulingIgnoredDuringExecution":
[{"weight":1, "preference": {"matchExpressions":
[{"key": "size", "operator": "In", "values":
["small"]}]}}]}}}'

Check placement
kubectl get pods -n learning -o wide
```

💡 Expected: nnappone2 runs on k3s-node2 (fallback!)

## 🧹 Clean Up

```
kubectl delete pod nnappone nnappone2 -n learning
kubectl taint node k3s-node1 test:NoSchedule- #
remove taint
kubectl label node k3s-node1 size- # remove label
kubectl delete namespace learning
```

# Required vs Preferred Affinity

| Feature                    | required...                                                                                         | preferred...                                                                                                               |
|----------------------------|-----------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------|
| <b>Constraint Type</b>     |  Hard (must match) |  Soft (try to match)                    |
| <b>Scheduling Failure?</b> |  Yes (if no match) |  No (always schedules)                  |
| <b>Use Case</b>            |  "Must run on GPU" |  "Prefer SSD, but OK on HDD"            |
| <b>Weight</b>              |  N/A             |  1-100 (higher = stronger preference) |



**Best Practice:** Use **preferred** for optimization, **required** for hard requirements

# Real-World Examples & Summary

## 💡 Cost Optimization Example

### \$ Weighted Preferences for Cost Savings

#### YAML Configuration

```
affinity: nodeAffinity:
 preferredDuringSchedulingIgnoredDuringExecution:
 - weight: 100 # Strong preference preference:
 matchExpressions: - key:
 node.kubernetes.io/instance-type operator: In
 values: [t3.small] # cheap spot instances
 - weight: 50 # Medium preference preference:
 matchExpressions: - key:
 topology.kubernetes.io/zone operator: In values:
 [us-east-1a]
```

"Strongly prefer cheap instances, and somewhat prefer zone A"

## 📄 Key Points

- ✓ **Soft constraint** — Never blocks scheduling
- ✓ **Weight (1-100)** — Express preference strength
- ✓ **Graceful fallback** — Works even if preferred nodes unavailable
- ✓ **Multi-node clusters** — Perfect for production environments
- ✓ **Combine with taints** — Simulate node unavailability

## 💡 Best Practice

Use **preferred** for optimization and **required** for hard requirements