# KUBERNETES NODE AFFINITY

## MULTIPLE VALUES & OR LOGIC: SCHEDULING OPTIMIZATION

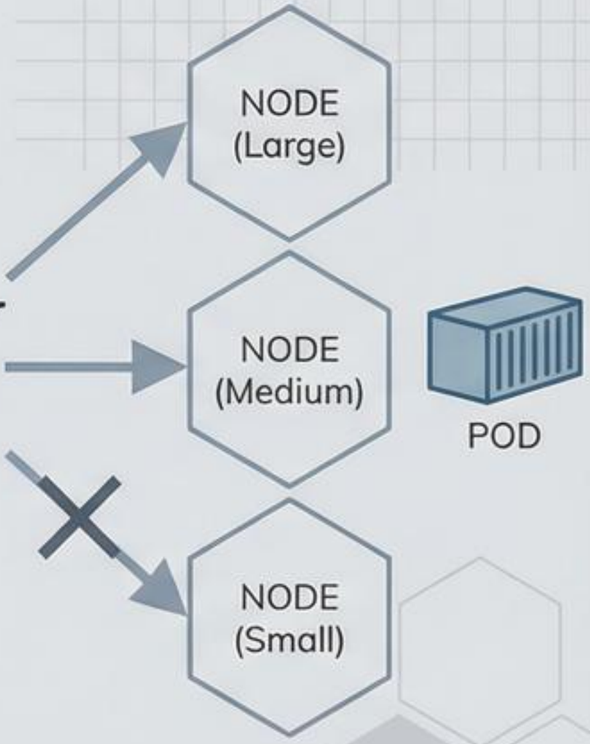



PRESENTATION DESIGNED FOR CONSULTING-LEVEL COMMUNICATION

# NODE AFFINITY: MULTIPLE VALUES & OR LOGIC

Understanding `size IN (large, medium)` in Kubernetes Scheduling

## YAML BREAKDOWN

```yaml
affinity:
  nodeAffinity:
    requiredDuringSchedulingIgnored
      DuringExecution:
      nodeSelectorTerms:
        - matchExpressions:
          - key: size
            operator: In
            values:
              - large
              - medium
```

🔑 **Key Insight:** Within one `matchExpressions`, `values: [large, medium]` translates to logical **OR**: `size IN (large, medium)`.

## LOGICAL EVALUATION PROCESS

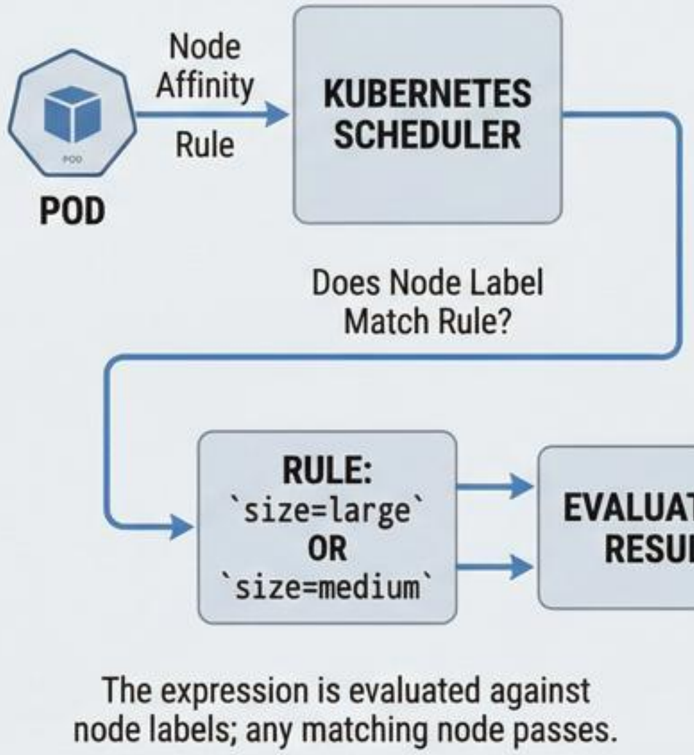


The expression is evaluated against node labels; any matching node passes.

## SCHEDULING OUTCOME

**NODE (Large)**
✅ Scheduled
(Matches `size=large`)

**NODE (Medium)**
✅ Scheduled
(Matches `size=medium`)

**NODE (Small)**
❌ Not Scheduled
(No Match)

The Pod can run on **ANY** node where the `size` label is either `large` OR `medium`. The filter is satisfied by either value.

# KUBERNETES NODE AFFINITY: TWO-LEVEL LOGIC STRUCTURE

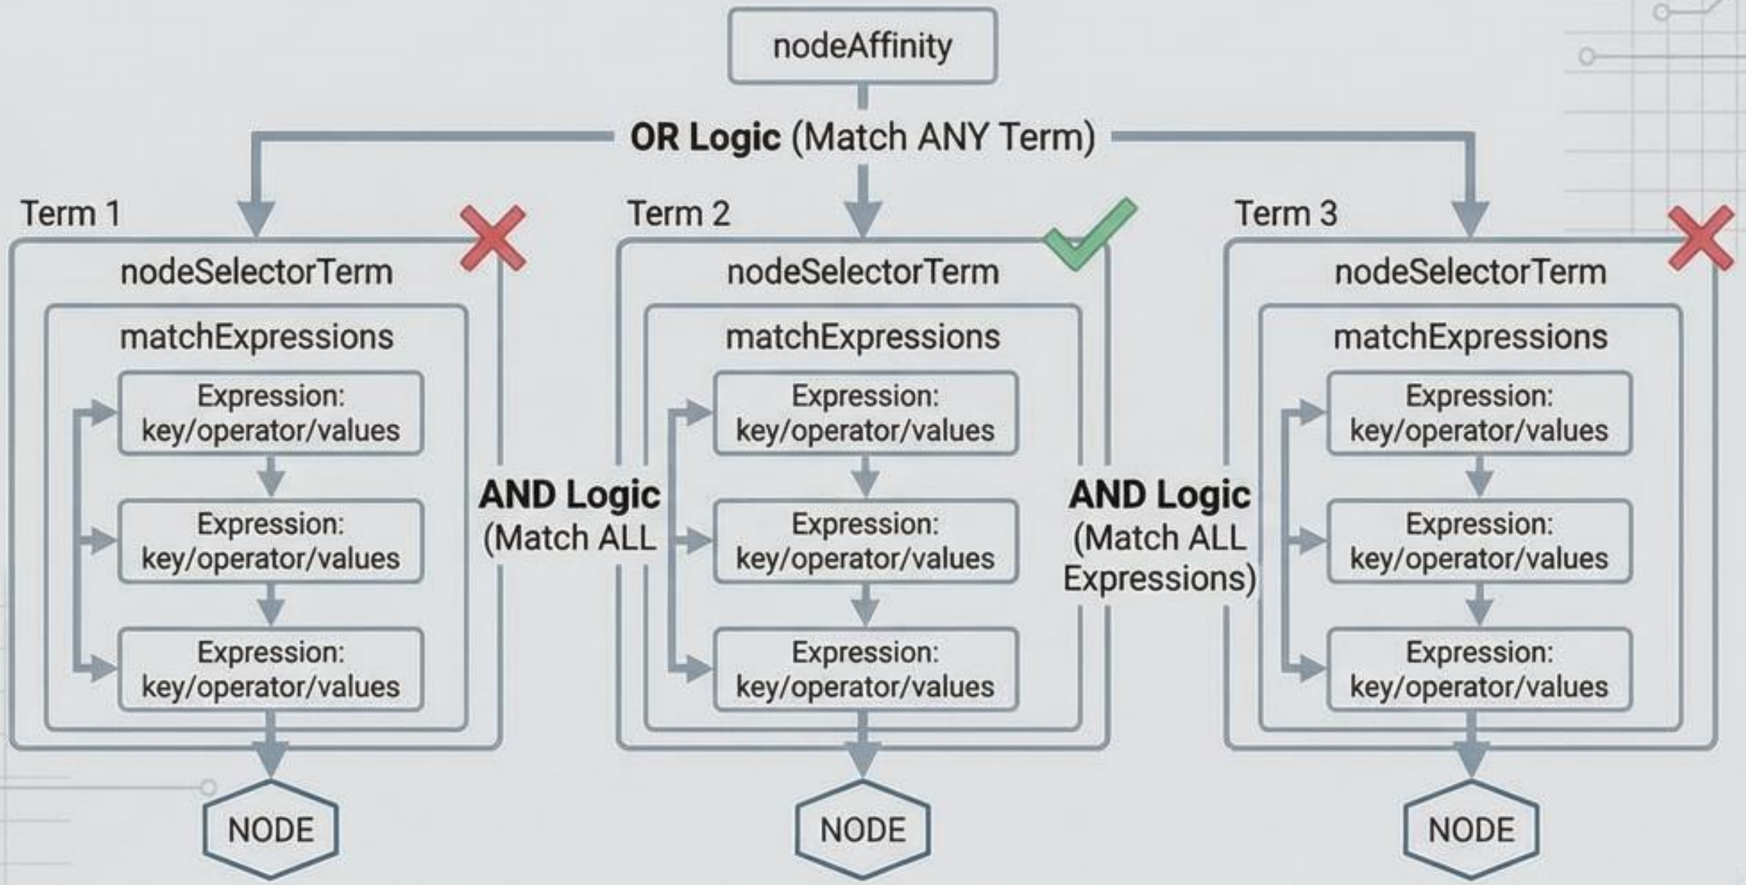


**Widen Choices:**

- Single Term, Multiple Values (e.g., size IN [large, medium])

**Narrow Choices:**

- Multiple Expressions in One Term (AND)
- Multiple Terms (OR different condition sets)

# K3s LAB: NODE LABELING & AFFINITY SETUP

## Prepare the Cluster for Multiple Value Testing (OR Logic)

### 1. LIST & LABEL NODES

```
$ kubectl get nodes
NAME          STATS  VERSION   AGE
k3s-master    eun    0         10m
k3s-node1     out    0         12m
k3s-node2     out    0         12m
```

```
kubectl label node
k3s-node1 size=small
```

```
kubectl label node
k3s-node2 size=large
```

Assign specific labels for testing.

### 2. VERIFY LABELS

```
$ kubectl get nodes --show-labels
NAME          LABELS
...
k3s-node1     ... size=small ✓
k3s-node2     ... size=large ✓
...
```


Confirm labels are successfully applied.

### 3. AFFINITY EVALUATION TEST BED

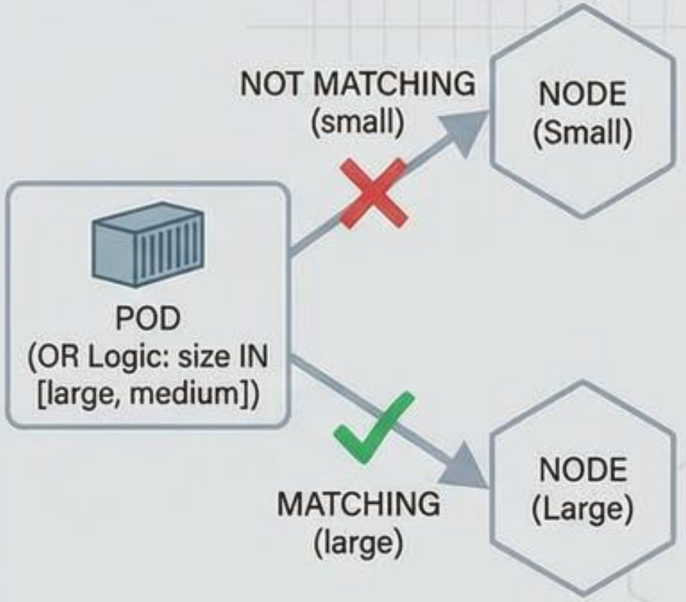


Labels create the test environment for immediate scheduling impact.

KEY INSIGHT: These labels form the foundation for testing OR logic; only nodes matching the Pod's multiple values are eligible for scheduling.

# KUBERNETES NODE AFFINITY: OR LOGIC IN ACTION

## Deploy Pod with Multiple Values & Verify Placement

**YAML**

```
# YAML Affinity
key: size
operator: In
values:
  - large
  - medium
```

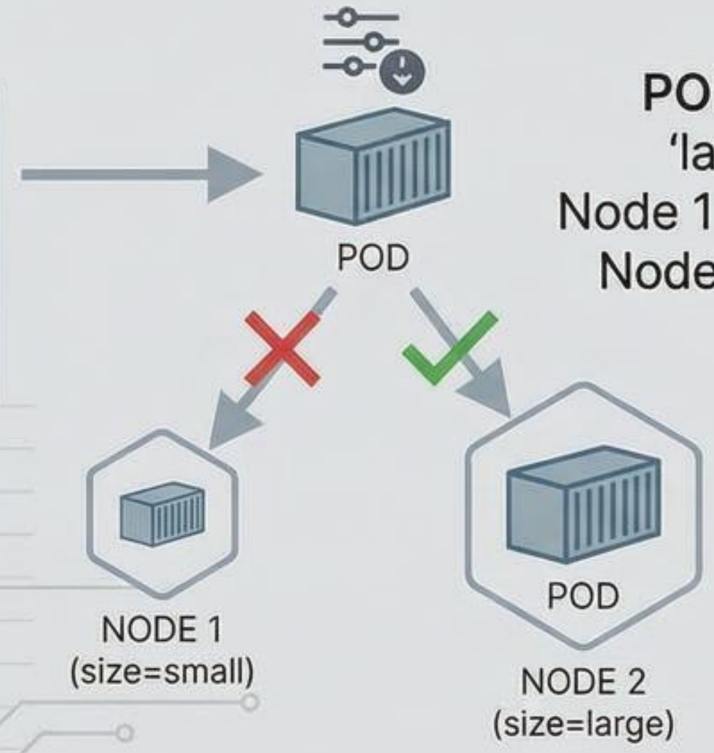


**POD AFFINITY** allows
'large' **OR** 'medium'.
Node 1 is 'small' (MISMATCH).
Node 2 is 'large' (MATCH).

```
# Check Pod Placement
# kubectl get pods -n learning -o wide
NAME      READY  STATUS   RESTARTS  AGE  IP          NODE
nnappone  1/1    Running  0         2m   10.42.0.10  k3s-node2
```

✅ POD LANDS ON
k3s-node2
(size=large)

Kubernetes multi-value list acts as an inclusive OR, selecting the first feasible node that meets at least one value, ignoring non-matching nodes.

# AFFINITY SCHEDULING FAILURE: NO MATCHING NODES (PENDING STATE)

## Demonstrating the Impact of Unmet Hard Affinity Rules (OR Logic)

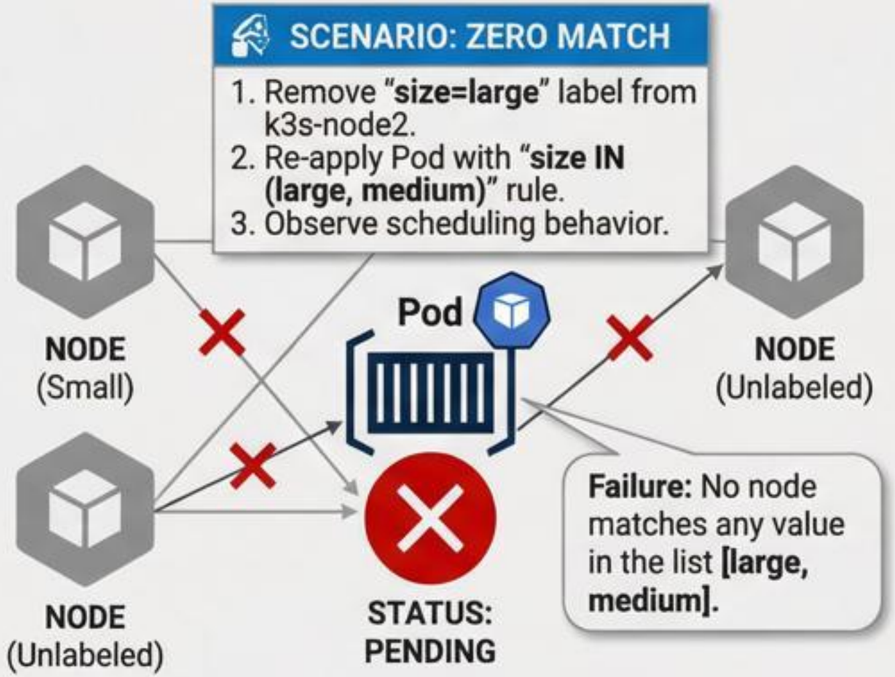


### k3s EXECUTION STEPS

```
# 1. Remove label from k3s-node2
kubectl label node k3s-node2 size-

# 2. Deploy Pod again
kubectl apply -f pod-with-node-affinity-multiple.yml

# 3. Check Pod status
kubectl get pods -n learning
```

### OBSERVED RESULTS

```
# Output from Step 3:
NAME       READY   STATUS    RESTARTS   AGE
nnappone   0/1     Pending   0          10s

# 4. Describe to confirm reason
kubectl describe pod nnappone -n learning

# Relevant Events Output:
Events:
  Type     Reason            Age From
  ----     ------            --- ----
  Warning  FailedScheduling  5s  default-scheduler 0/2 nodes match node affinity: node(s) didn't match node selector.
```

### KEY INSIGHT: OR LOGIC REQUIRES AT LEAST ONE MATCH

The exercise proves that OR logic does not relax hard affinity. If no node satisfies any value in the list, scheduling blocks indefinitely until a qualifying node appears or the rule is changed. Hard affinity is a strict filter.

# CLUSTER BASELINE: RETURNING TO A CLEAN SLATE

Remove lab artifacts to ensure consistent future environments.

POD

## 1. DELETE TEST POD

```
kubectl delete pod nnappone -n learning
```

Removes the application workload.

## 2. STRIP SIZE LABELS

```
kubectl label node k3s-node1 size-
kubectl label node k3s-node2 size-
```

Prevents stale labels from influencing placements.

## 3. REMOVE NAMESPACE

```
kubectl delete namespace learning
```

Eliminates all associated resources.

**PRO TIP:** Consistent cleanup habits avoid surprise affinity matches during later demos and keep GitOps diffs minimal.

# KUBERNETES NODE AFFINITY: MULTIPLE VALUES & OR LOGIC (REAL-WORLD USE CASES)

## SCHEDULING OPTIMIZATION: AVOIDING RESOURCE STARVATION

### PRODUCTION CLUSTERS: CAPACITY & LIFECYCLE

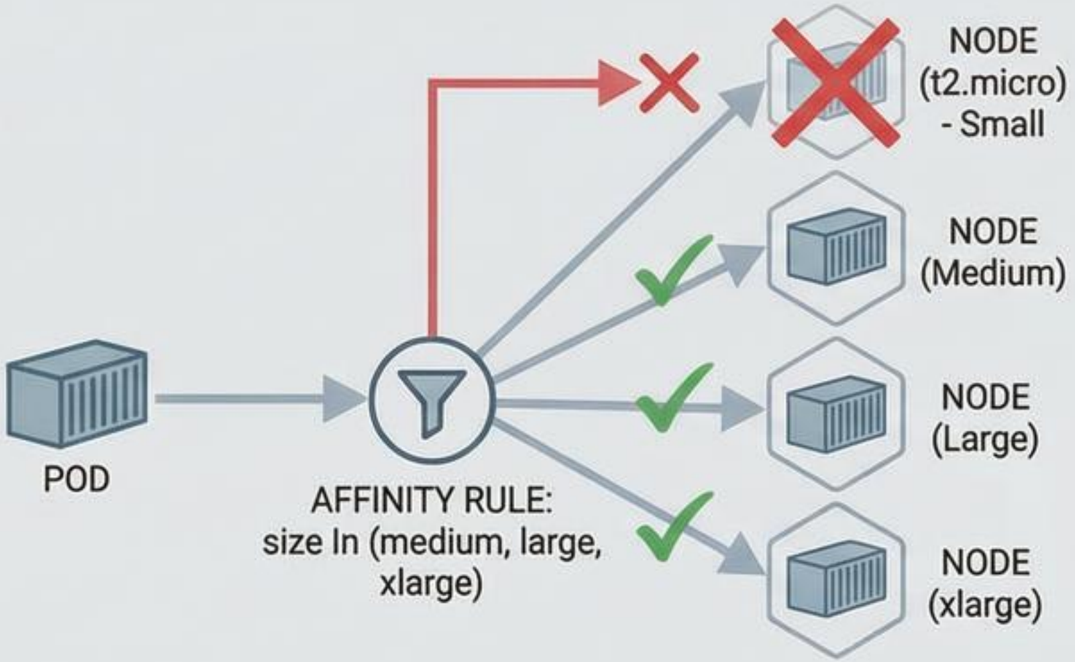


Avoid Small Instances. Widens placement pool to stable sizes while excluding t2.micro equivalents.

### AFFINITY LOGIC: SINGLE EXPRESSION, MULTIPLE TIERS

```
affinity:
  nodeAffinity:
    requiredDuringSchedulingIgnoredDuringExecution:
      nodeSelectorTerms:
      - matchExpressions:
          - key: size
            operator: In
            values:
              - medium
              - large
              - xlarge
```

- **PREVENTS RESOURCE STARVATION**: Ensures Pods have adequate resources, reducing eviction risk.
- **COMBINES ACCEPTABLE TIERS**: Illustrates how OR logic simplifies rules by grouping multiple acceptable values.
- **REDUCES EVICTION RISK**: Prioritizes stable, sufficient nodes over potentially unstable small instances.

# Node Affinity Logic: Combining OR and AND within a Single Term

## Real-World Use Case: Multi-AZ, Multi-GPU Pod Scheduling

Goal: Spread across zones (OR) AND restrict to specific GPUs (AND)

```
nodeSelectorTerms:
- matchExpressions:
  - key: topology.kubernetes.io/zone
    operator: In
    values:
      - us-east-1a
      - us-east-1b
  - key: accelerator
    operator: In
    values:
      - nvidia-tesla-v100
      - nvidia-tesla-t4
```

OR Logic (Fault Tolerance)

OR Logic (Hardware Choice)

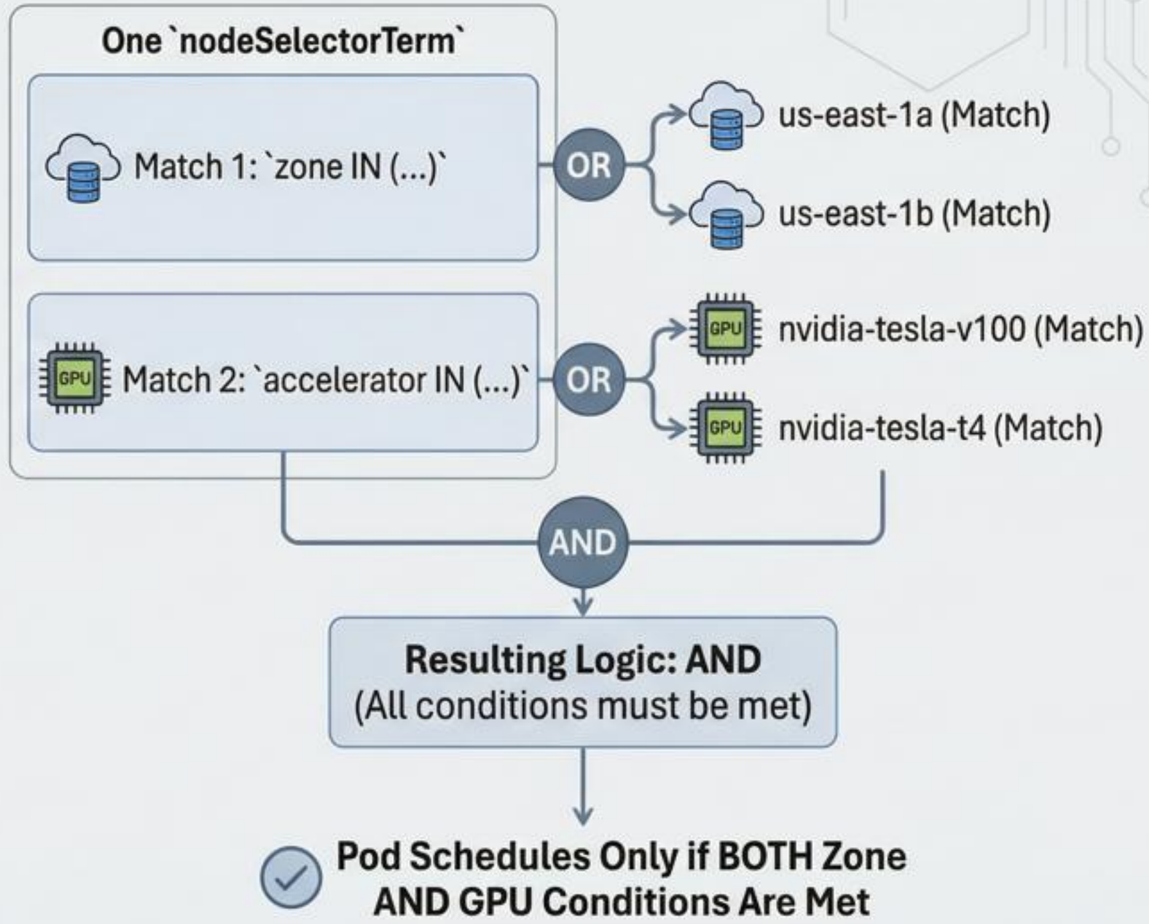

# KUBERNETES NODE AFFINITY: MULTIPLE VALUES & LOGIC OPTIMIZATION

Addressing Common Questions: OR, AND, Negation, and Combination.

## OR Logic: `values: [large, medium]` (Same `matchExpression`)

```
matchExpressions:
 - key: size
   operator: In
   values:
    - large
    - medium
```

**size IN** (large, medium) → **Logical OR**

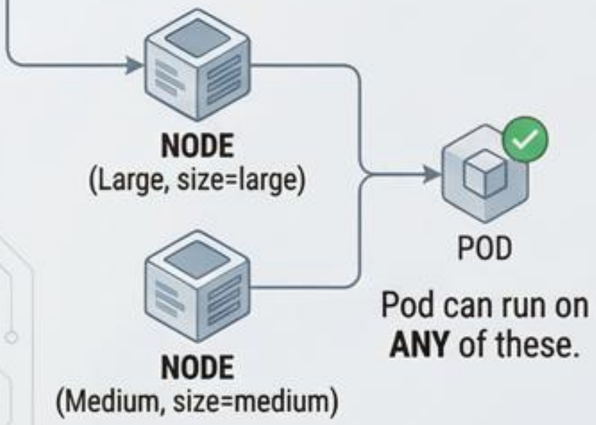


Pod can run on **ANY** of these.

Key Insight: Single `matchExpression` with multiple values acts as OR.

## AND Logic: Multiple `matchExpressions` (Same Term)

```
nodeSelectorTerms:
 - matchExpressions:
   - key: size
     operator: In
     values: [large]
   - key: disk
     operator: In
     values: [ssd]
```

**size=large AND disk=ssd** → **Logical AND**

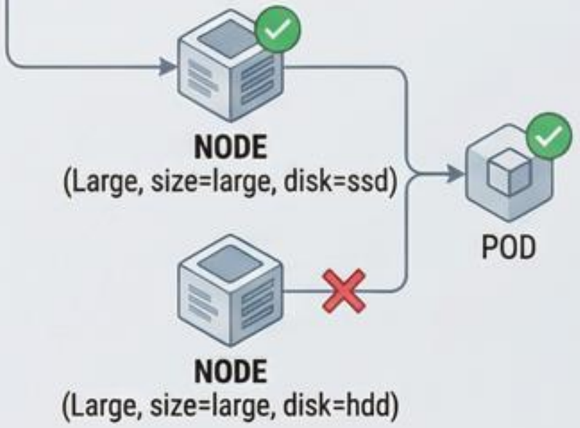


Pod requires ALL conditions to be true.

## Negation: `operator: NotIn`

```
matchExpressions:
 - key: size
   operator: NotIn
   values:
    - small
    - medium
```

**size NOT IN** (small, medium)

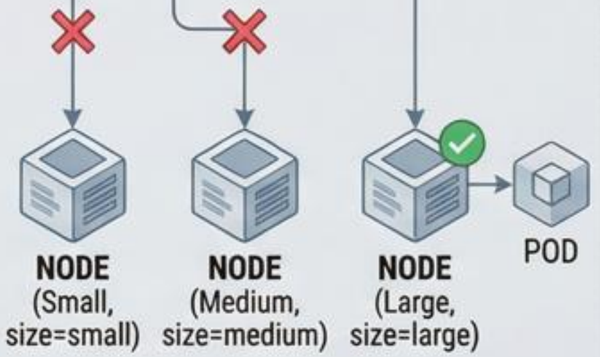


Explicitly excludes unwanted nodes (e.g., small, medium). Equivalent candidate set to In [large, xlarge] if only those exist.

## Practical Tips

**Value Order Irrelevant:**
`values: [A, B]` is identical to `values: [B, A]`.

`[large, medium]` = `[medium, large]`

**Combining Required & Preferred:**
Use both for fine-grained control.

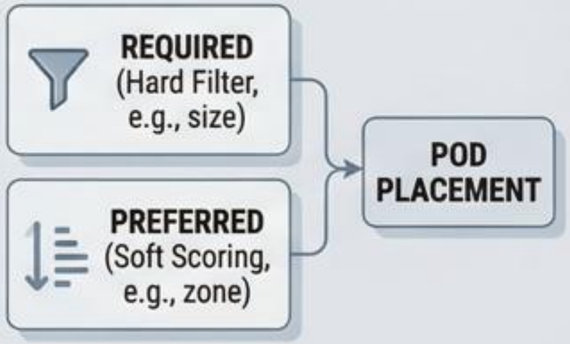


Implement hard constraints plus soft scoring for optimization.

PRESENTATION DESIGNED FOR CONSULTING-LEVEL COMMUNICATION

YAML

# Node Affinity Logic: The Three-Rule Mantra

## 1. Within a Single matchExpression

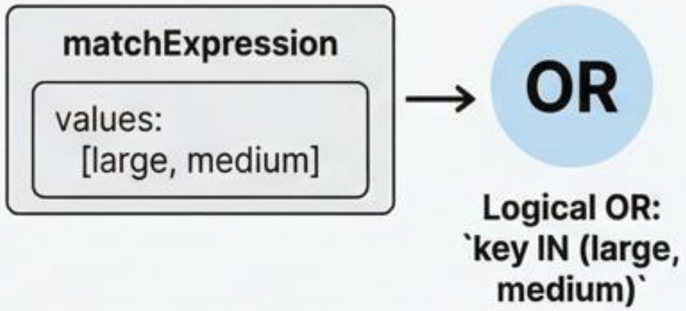


Use for OR on one label.

## 2. Across Multiple matchExpressions

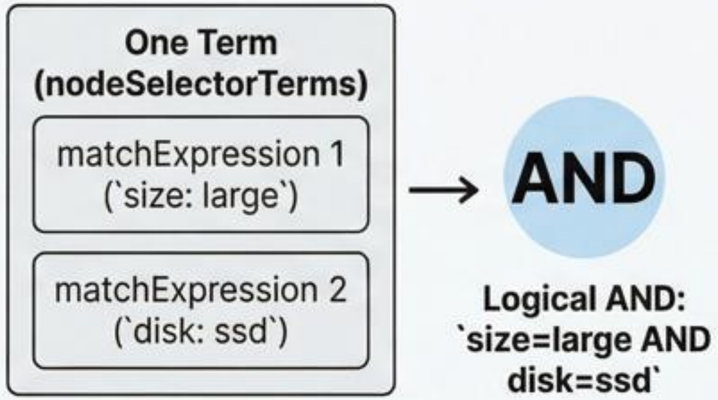


Use for AND across labels.

## 3. Across Multiple nodeSelectorTerms

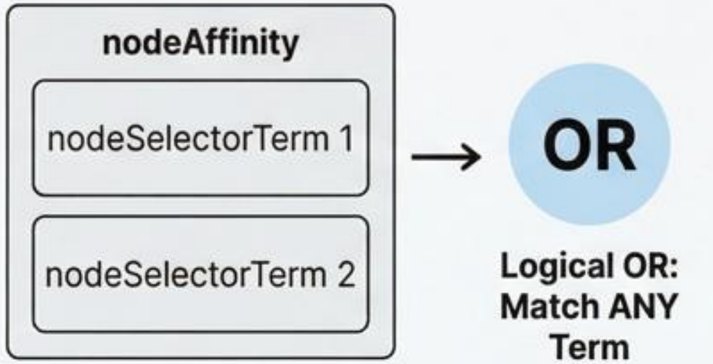


Use for OR between groups of conditions.

## Mastery & Benefits

- Read manifests quickly.
- Predict scheduling outcomes.
- Keep clusters reliable & intentions clear.