

List of possible errors of the Kubernetes ecosystem

Resolving issues in Kubernetes can be systematic and efficient if you follow a structured approach. So, in this article, we will try to discuss different objects and associated possible errors and we will also develop a progressive and right directional approach to figure out core reason of causes.

1. Pod and Container Issues

- **CrashLoopBackOff**: Pod is crashing repeatedly due to incorrect configurations or application failures.
- **ImagePullBackOff**: Kubernetes cannot pull the specified image from the container registry.
- **ErrImagePull**: Error pulling the image (due to incorrect image name, tag, or authentication issues).
- **OOMKilled**: Pod is terminated because it exceeds its memory limits.
- **ContainerCreating**: Pod is stuck in ContainerCreating state due to volume, network, or image pull issues.
- **Completed Pods**: Pods are stuck in Completed state due to job completion but not being cleaned up.
- **Terminating Pods**: Pods are stuck in Terminating state, often caused by finalizer misconfigurations or unresponsive cleanup.
- **PodEviction**: Pods are evicted from nodes due to resource pressure or taints and tolerations mismatch.

2. Deployment and ReplicaSet Issues

- **ReplicaSet Misalignment**: Desired replicas don't match actual replicas due to scheduling failures or insufficient resources.
- **FailedScheduling**: Pods fail to schedule due to resource constraints, taints, or node affinity/anti-affinity rules.
- **Rolling Update Issues**: New deployments cause downtime due to misconfigured readiness/liveness probes or missing maxUnavailable/maxSurge settings.
- **HPA (Horizontal Pod Autoscaler) Issues**:
 - HPA not scaling due to missing metrics (e.g., CPU/Memory not exposed).
 - Incorrect resource requests/limits affecting scaling decisions.

3. Networking Issues

- **DNS Resolution Failures**: Pods cannot resolve service names due to CoreDNS misconfigurations or failures.
- **Connection Refused**: Services are not accessible due to misconfigured or absent network policies.
- **Service Not Reachable**: Issues with ClusterIP, NodePort, or LoadBalancer services caused by configuration errors or cloud provider limitations.
- **NetworkPolicy Blocking Traffic**: Improperly defined NetworkPolicies block desired ingress/egress traffic.

- **Kube-proxy Issues:** Problems with service networking due to kube-proxy misconfigurations.

4. Storage Issues

- **VolumeMount Failure:** Persistent Volumes (PVs) fail to attach or mount to Pods.
- **PVC Pending:** Persistent Volume Claims (PVCs) remain in a Pending state due to insufficient or unavailable storage.
- **Data Loss:** Misconfigured StatefulSets or volume binding lead to data loss when Pods are rescheduled.

5. Control Plane Issues

- **API Server Unreachable:** Issues due to high load, misconfiguration, or certificate problems.
- **etcd Performance Degradation:** etcd cluster has high latency or data corruption, affecting the overall cluster performance.
- **Controller Manager Issues:** Failures in applying desired state due to misconfigurations or bugs.
- **Scheduler Failures:** Scheduler fails to allocate Pods to nodes due to configuration or resource exhaustion.

6. Authentication and Authorization Issues

- **RBAC Denied Errors:** Users or Pods cannot perform actions due to missing or misconfigured Role-Based Access Control (RBAC) permissions.
- **Service Account Issues:** Missing or misconfigured Service Accounts cause authentication failures for workloads.
- **Webhook Authentication Errors:** External authenticators fail due to connectivity or misconfigurations.

7. Node Issues

- **Node Not Ready:** Nodes are marked NotReady due to kubelet, disk pressure, or network issues.
- **DiskPressure or MemoryPressure:** Nodes face resource exhaustion, evicting Pods or causing degraded performance.
- **Uncordoned Nodes:** Nodes remain cordoned after maintenance or upgrades, causing scheduling failures.
- **Kubelet Failures:** Kubelet crashes or becomes unresponsive, affecting Pods on the node.

8. Ingress and Load Balancer Issues

- **Ingress Not Routing:** Misconfigured Ingress resources or absent Ingress controllers.
- **SSL/TLS Errors:** Missing or improperly configured certificates for HTTPS traffic.
- **LoadBalancer Not Created:** Cloud Load Balancer fails to provision due to provider-specific limitations or misconfigurations.
- **Backend Service Unhealthy:** Services backing the Ingress are not passing health checks.

9. Configuration Issues

- **Environment Variables Missing:** Pods crash due to missing or incorrectly set environment variables.
- **Secret and ConfigMap Issues:** Secrets or ConfigMaps are not mounted correctly, causing application errors.
- **Improper Resource Requests and Limits:** Over or under-provisioned resource requests causing scheduling or runtime failures.

10. Logging and Monitoring Issues

- **Log Aggregation Issues:** Centralized logging tools (e.g., Fluentd, Logstash) fail to collect or forward logs.
- **Metrics Missing:** Metrics server or Prometheus is misconfigured, leading to missing metrics for scaling or monitoring.
- **Application Debugging Issues:** Lack of application-specific logs due to inadequate logging configurations.

11. Cluster Upgrades and Maintenance Issues

- **Version Skew:** Components (e.g., kube-apiserver, kubelet) running incompatible versions after upgrades.
- **Downtime During Upgrade:** Misconfigured workloads cause downtime during control plane or node upgrades.
- **Custom CRD Breakages:** Breaking changes in custom resource definitions (CRDs) due to API deprecations.

12. Scaling and Performance Issues

- **Node Resource Exhaustion:** Nodes run out of CPU, memory, or disk, causing degraded performance or Pod evictions.
- **Cluster Autoscaler Failures:** Cluster Autoscaler fails to provision nodes due to cloud provider limitations or incorrect configurations.
- **Throttling in API Server:** High API request volume causes API server throttling.

13. Debugging and Observability Challenges

- **Lack of Metrics or Logs:** Missing metrics/logs make debugging difficult.
- **Event Overwrites:** Kubernetes events have a short retention window, leading to lost debugging context.
- **Misinterpreting Liveness/Readiness Probes:** Incorrectly configured probes cause unnecessary Pod restarts.